

NØNOS Bootloader Architecture

UEFI Secure Boot • Ed25519 Signatures • Groth16 ZK Proofs • 185 Rust Files

nonos-boot/src/

loader/	crypto/	zk/	entropy/	security/	network/	handoff/	verify/
ELF64 Parser Segment Loading Relocation (PIE) Memory Allocation 42 files	Ed25519 Verify BLAKE3 Hashing Production Keys KeyStore (16 max) 4 files	Groth16 Verify BLS12-381 Pairing Circuit Registry Boot Attestation 14 files	RDRAND/RDSEED TSC Jitter UEFI RNG Protocol RTC Entropy 6 files	Secure Boot Check TPM PCR Extend Platform Security Crypto Self-Test 7 files	PXE Boot TFTP Download HTTP Fetch DHCP/Static IP 15 files	BootHandoffV1 Memory Map Copy ExitBootServices Kernel Jump 6 files	Capsule Validation Signature Check Hash Verification Integrity Chain 4 files
boot/	config/	hardware/	capsule/	multiboot/	secure_loader/	crypto_visual/	testing/
Logo Display Phase Banners Status Output 9 files	NVRAM Settings Boot Policies Configuration 7 files	CPU Detection ACPI Tables Memory Info 7 files	Capsule Format ZK Metadata Kernel Extract 3 files	NVRAM Prefs Boot Manager Entry Selection 4 files	Secure Loading Measurements Validation 14 files	BLAKE3 Display Ed25519 Visual ZK Status 13 files	Test Framework Memory Tests Crypto Tests 12 files

Boot Flow



185 Rust Files | 20 Subsystems | 3 Ed25519 Keys
3 ZK Circuits | 7 Domain Strings | 440 NONOS Refs

Magic: 0x4E4F4E4F ("NONO")
Handoff: BootHandoffV1 + 32KB Stack

Target: NØNOS Kernel
extern "C" fn(BootHandoffV1*)

Core Subsystem | Support Module
UEFI/Kernel Interface

UEFI Firmware



NØNOS Boot Flow

UEFI Entry → Security Init → Kernel Load → Verify → Handoff → Jump



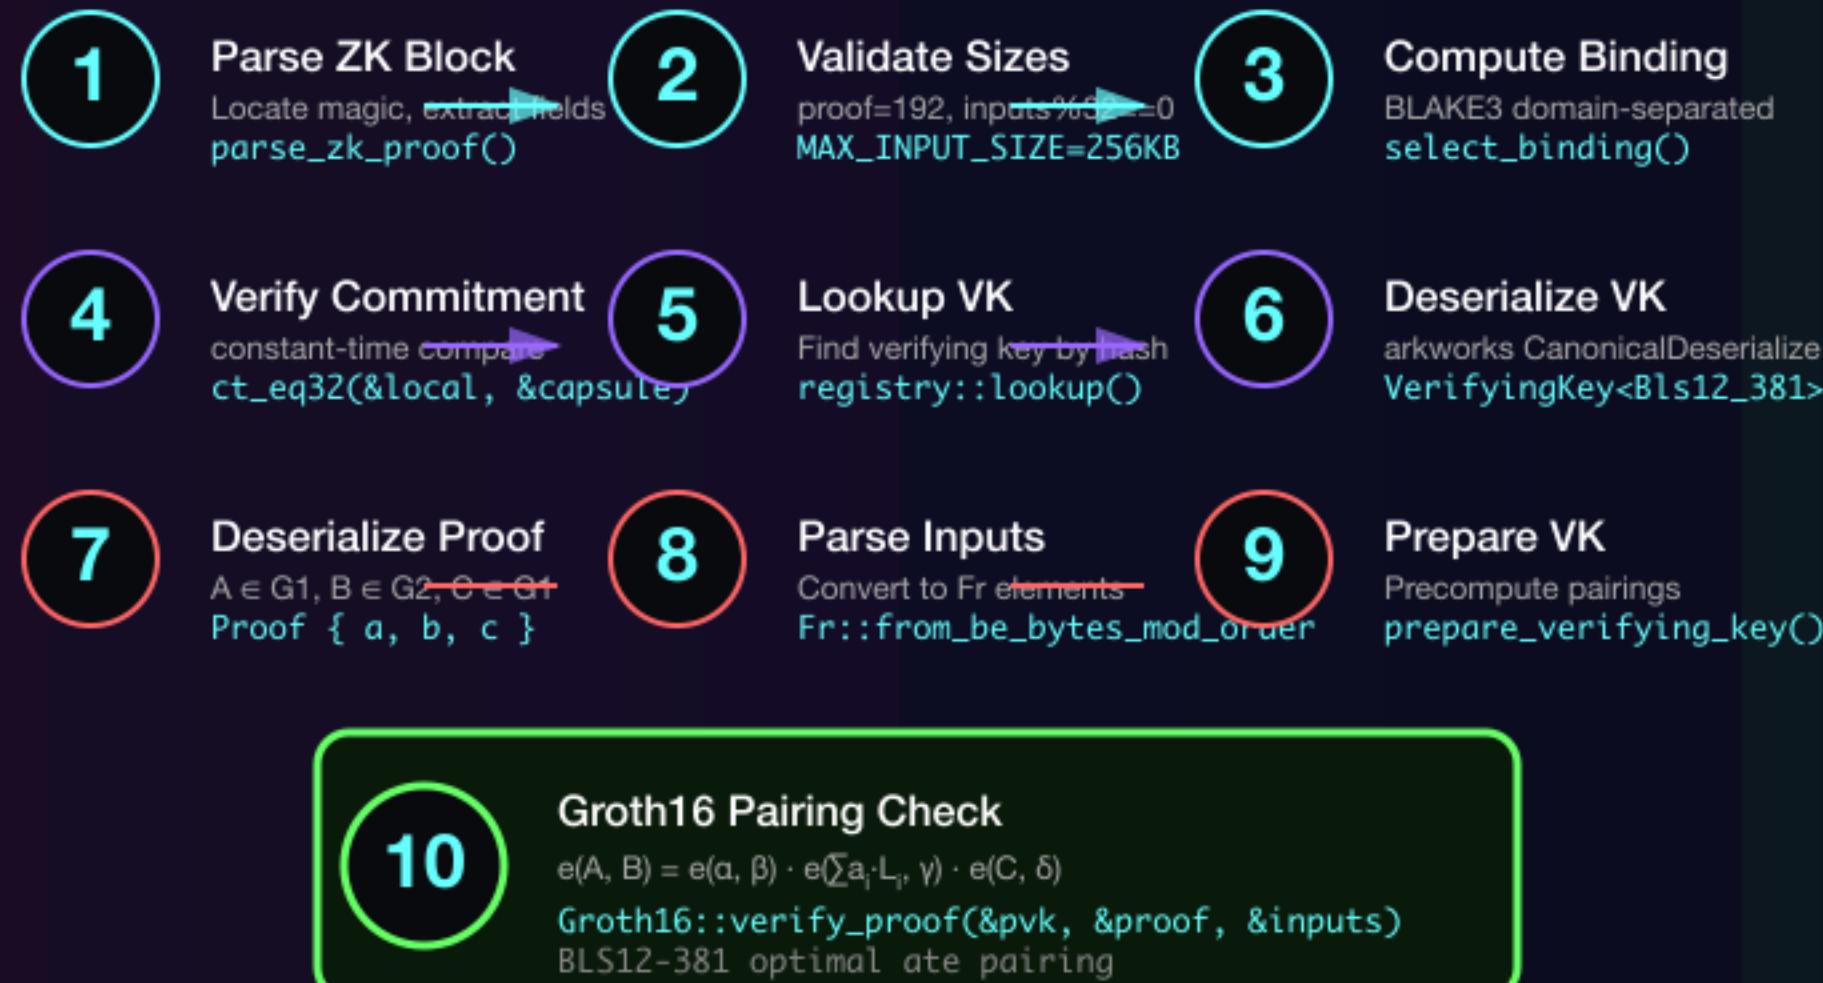
ZK Boot Attestation

Groth16 ZK-SNARK Verification • BLS12-381 Pairing • Zero-Knowledge Boot Proof

ZK Proof Block Format

Magic	[4 bytes]	[0x4E, 0xC3, 0x5A, 0x50]	"NoZP"
Version	[4 bytes]	1 (little-endian u32)	
Program Hash	[32 bytes]	BLAKE3 circuit identifier	
Capsule Commitment	[32 bytes]	BLAKE3 binding hash	
Public Inputs Len	[4 bytes]	$n * 32$ (field elements)	
Proof Blob Len	[4 bytes]	192 (Groth16 compressed)	
Public Inputs	[n bytes]	Fr field elements 32-byte aligned	
Groth16 Proof	[192 bytes]	$A (G1) + B (G2) + C (G1)$ $48 + 96 + 48 = 192$ bytes compressed	

Verification Pipeline



Circuit Registry

boot-authority	Kernel boot authorization <code>BootAuthority</code> <code>Attestation</code>
update-authority	System update authorization <code>UpdateAuthority</code>
recovery-key	Recovery operation auth <code>RecoveryKey</code>
Verification Key	568+ bytes (BLS12-381) Compressed canonical form

ENTRIES: $\{ (\text{program_hash}, \text{vk_bytes}) \}$

Domain Separation Strings

Program Hash:	"NONOS:ZK:PROGRAM:v1"
Commitment:	"NONOS:CAPSULE:COMMITMENT:v1"
Circuit Key:	"NONOS:CIRCUIT_KEY:v1"

All use `BLAKE3::new_derive_key()` for domain separation

Verification Result

ZkVerifyResult::Valid	Proof verified successfully
ZkVerifyResult::Invalid	Proof or commitment failed
ZkVerifyResult::Unsupported	Unknown program hash
ZkVerifyResult::Error	Deserialization or internal

Passed to Kernel (BootHandoffV1.zk)

ZkAttestation Struct	BootAttestationResult
verified: u8 (1 = valid) flags: u8 (reserved) program_hash: [u8; 32] capsule_commitment: [u8; 32]	zk_verified: bool program_hash: [u8; 32] capsule_commitment: [u8; 32] detail: ZkVerifyResult

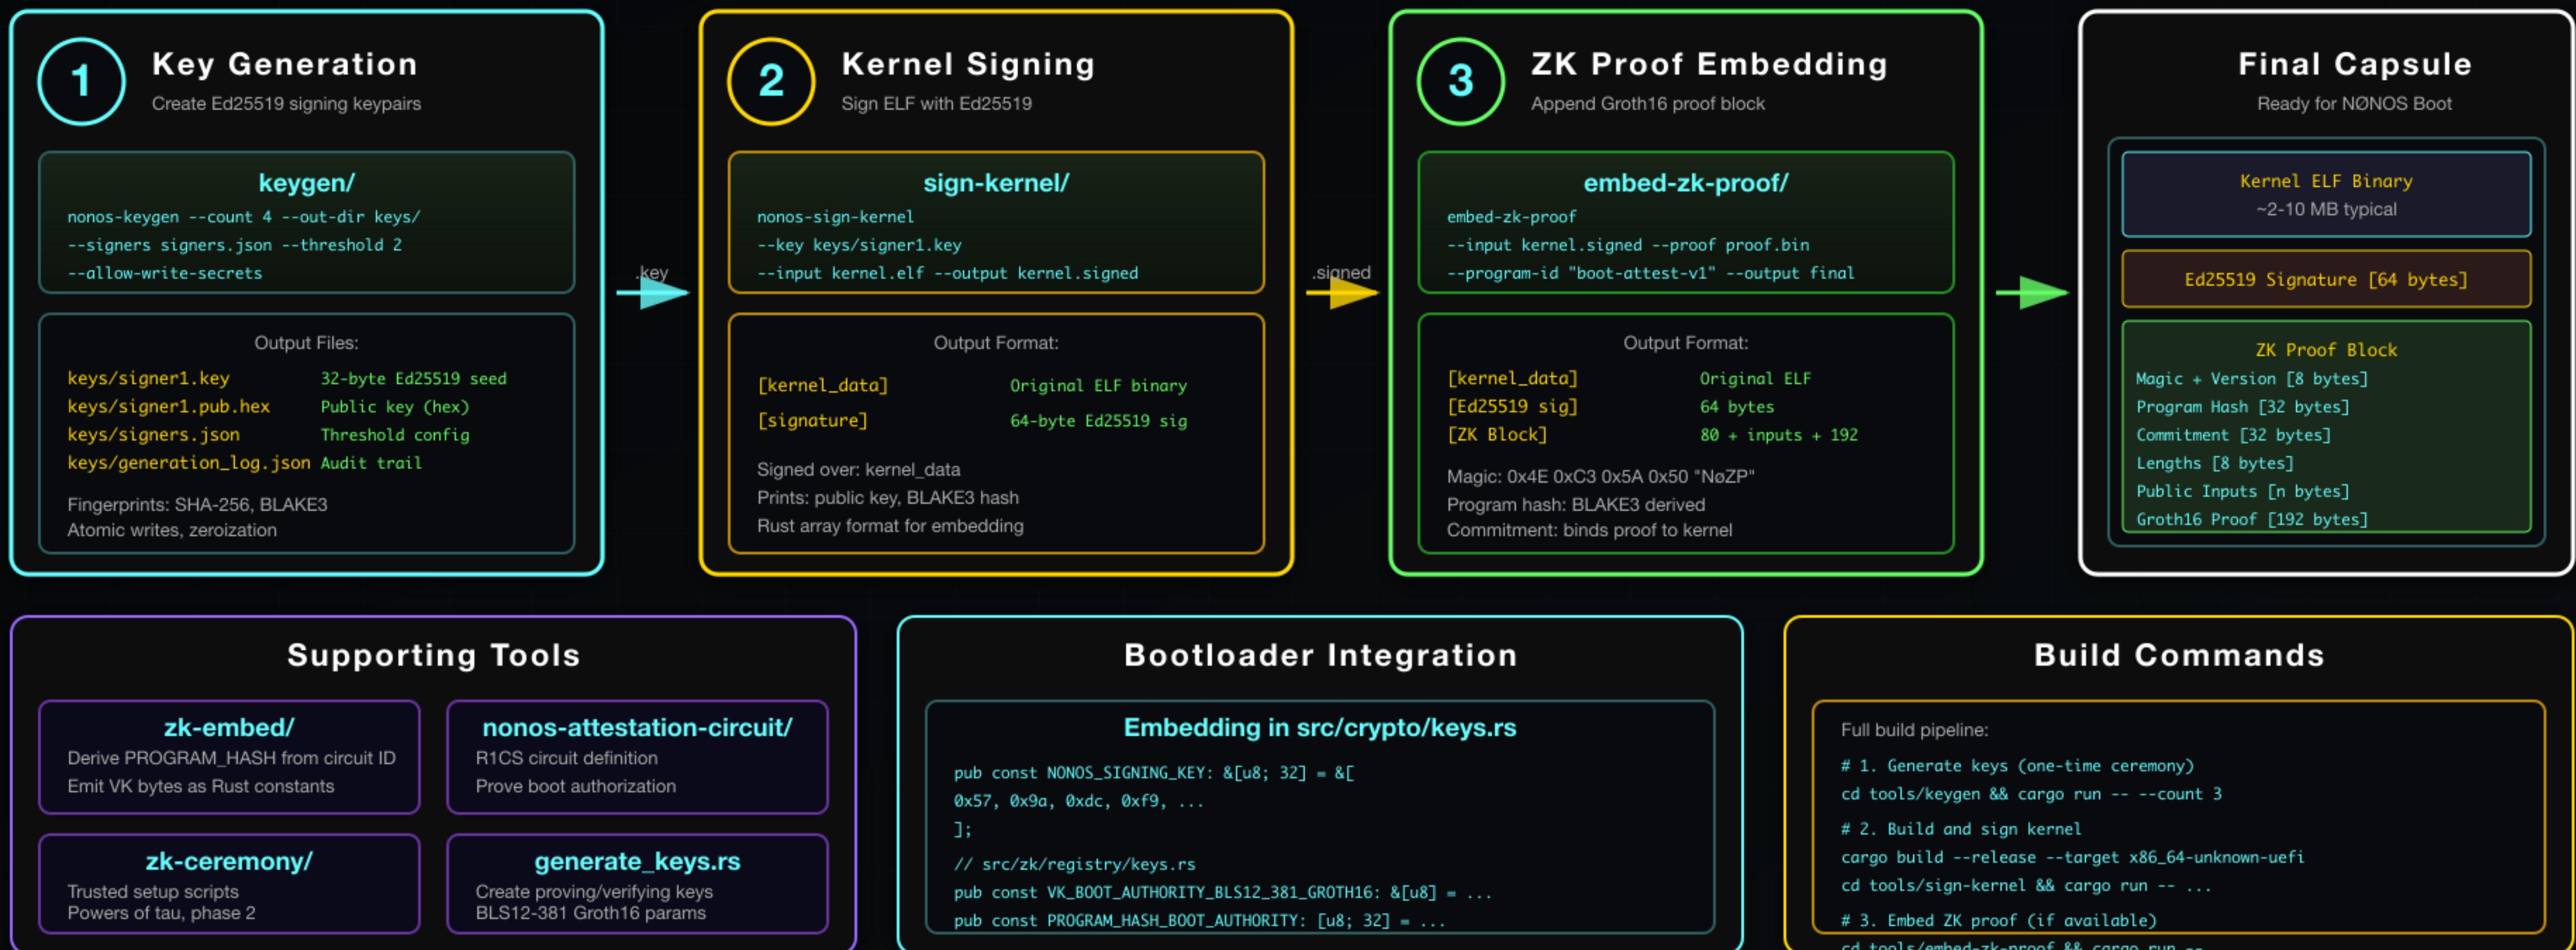
Security Properties

- ✓ Zero-Knowledge
- ✓ Soundness
- ✓ Binding
- ✓ Non-malleable
- ✓ Constant-time
- ✓ Zeroization

Proof reveals nothing about witness
Cannot forge proof without witness
Commitment ties proof to kernel
Proof cannot be modified
Side-channel resistant comparison
Proof wiped after verification

NØNOS Build Tools

nonos-boot/tools/ • Key Generation → Kernel Signing → ZK Proof Embedding



NONOS Bootloader: Cryptographic Subsystem

Ed25519 Signatures + BLAKE3 Hashing + Domain Separation

KEYSTORE

src/crypto/keys.rs

KeyStore Structure
keys: [[u8; 32]; 16] // MAX_KEYS
count: usize
static KEYSTORE: Mutex<KeyStore>
static INIT_DONE: AtomicBool

Production Keys (Ed25519)

Key 1 (Primary):
579adc...9b7ffb (NONOS_SIGNING_KEY)
Key 2 (Backup):
3d4017...f4660c
Key 3 (Recovery):
fc51cd...908025

Embedded at compile time

BLAKE3 HASHING

Domain Separation Strings

KeyID Derivation:

"NONOS:KEYID:ED25519:v1"
derive_keyid(pubkey) -> [u8; 32]

Entropy Accumulation:

"NONOS:BOOT:ENTROPY:ACCUM"
collect_boot_entropy() mixing

Entropy Output:

"NONOS:BOOT:ENTROPY:OUTPUT"
Final entropy extraction (XOF)

ZK Program Hash:

"NONOS:ZK:PROGRAM:v1"
Circuit VK normalization

ZK Commitment:

"NONOS:CAPSULE:COMMITMENT:v1"

SIGNATURE VERIFICATION

src/crypto/verify.rs

verify_signature_bytes()

1. Check sig.len() == 64 (SIG_LEN)
2. Check is_initialized()
3. Parse Ed25519 Signature
4. For each key in KEYSTORE:
 - VerifyingKey::from_bytes()
 - pk.verify(data, &sig)
5. Return KeyId on success

verify_signature_full()

1. Extract sig from blob[meta.offset]
2. Extract payload from blob
3. Reject all-zero signatures
4. Call verify_signature_bytes()

Uses ed25519-dalek crate

CRYPTO SELF-TEST

perform_crypto_self_test()

BLAKE3 Test:

hash("NONOS-crypto-selftest") x2
Verify h1 == h2 (determinism)

Ed25519 Test:

VerifyingKey::from_bytes(NONOS_KEY)

VERIFY ERRORS

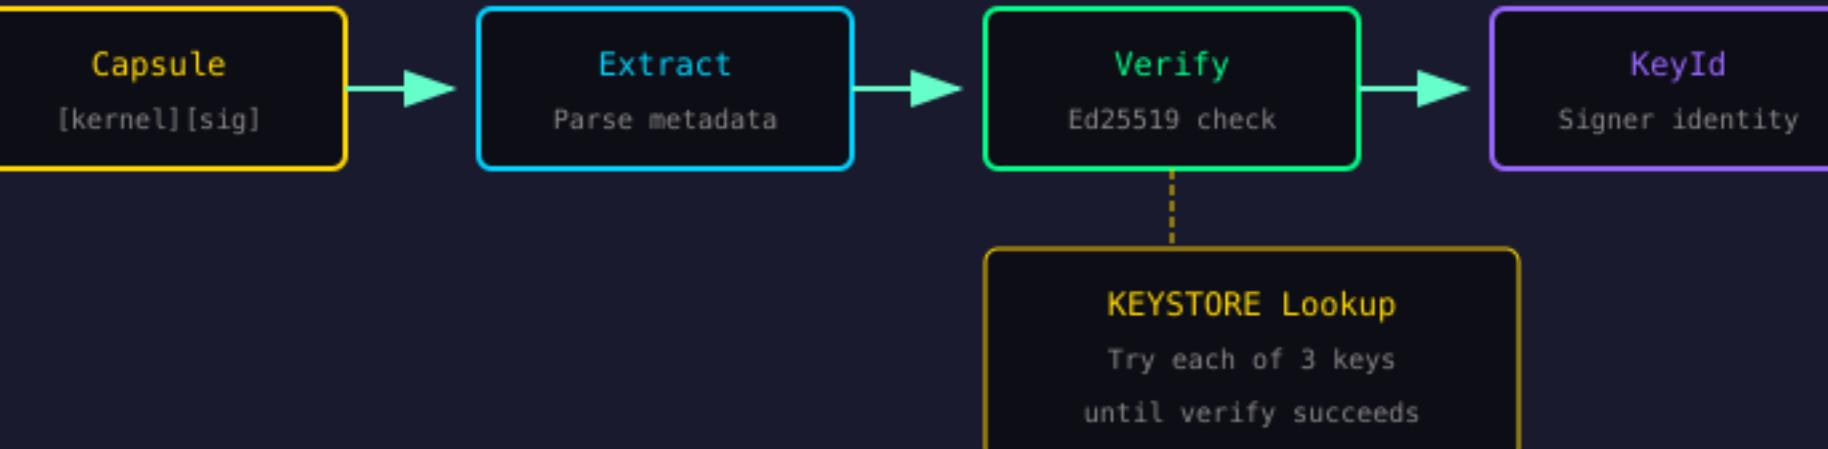
VerifyError enum:

Bounds – offset/length overflow
MalformedSignature – bad format
InvalidSignature – verify failed
KeyNotFound – no matching key
NotInitialized – keystore empty

CAPSULE METADATA

CapsuleMetadata struct:
offset_sig: usize // Signature location
len_sig: usize // 64 bytes (SIG_LEN)
offset_payload: usize // Kernel data start
len_payload: usize // Kernel size
signer_keyid: Option<KeyId>
payload_hash: [u8; 32] // BLAKE3 hash
header_version: u32, header_timestamp: u64

SIGNATURE VERIFICATION FLOW



CRYPTO STATISTICS

Algorithms:

Ed25519 – 256-bit elliptic curve signatures
BLAKE3 – Fast cryptographic hash (XOF)
SHA-256 – TPM PCR measurements

Key Properties:

Public Key: 32 bytes (PK_LEN)
Signature: 64 bytes (SIG_LEN)
KeyId: 32 bytes (BLAKE3 derived)

Cryptographic Foundation: Zero external network calls | All keys embedded at compile time | Domain-separated hashing prevents cross-protocol attacks

Dependencies: ed25519-dalek, blake3, sha2 | Crates verified against known hashes

NONOS Bootloader: Entropy Collection

Multi-Source Hardware Entropy for Kernel RNG Seed

EFI_RNG PROTOCOL

```
#[cfg(feature = "efi-rng")]
```

```
collect_efi_rng()
1. bs.locate_protocol::<Rng>()
2. rng.get_rng(None, &mut buf[64])
3. h.update(&buf)
4. scrub(&mut buf) // Zero after use

UEFI firmware hardware RNG
```

RDRAND / RDSEED

Intel/AMD Hardware RNG Instructions

```
collect_hw_rng_bytes()
HW_RNG_ITERATIONS = 32
For each iteration:
    - Try RDSEED first (true entropy)
    - Fallback to RDRAND (PRNG)
    - Accumulate into 64-byte buffer
CPU instruction-level entropy
```

TSC JITTER

Timing Variations from CPU

```
collect_tsc_jitter()
TSC_JITTER_ROUNDS = 256
For each round:
    - t1 = rdtsc_serialized()
    - bs.stall(23 + (round*7 % 17))
    - t2 = rdtsc_serialized()
    - delta = t2 - t1 (with rotation)
Unpredictable timing variations
```

RTC ENTROPY

Real-Time Clock Data

```
collect_rtc_entropy()
st.runtime_services().get_time()
16-byte buffer contains:
    - year (4B) + month + day
    - hour + minute + second
    - nanosecond (4B) + timezone
    - daylight savings flag
Nanosecond precision adds entropy
```

BLAKE3 ENTROPY MIXING

Domain-Separated Accumulation and Output Extraction

Accumulation Phase

```
"NONOS:BOOT:ENTROPY:ACCUM"
blake3::Hasher::new_derive_key(DS_ENTROPY_ACCUM)
```

Output Phase (XOF)

```
"NONOS:BOOT:ENTROPY:OUTPUT"
finalize_xof().fill(&mut out[64]) // 64 bytes
```

WEAK ENTROPY RESCUE

Fallback when entropy quality is low

```
rescue_weak_entropy()
1. rescue = rdtsc_serialized()
2. BLAKE3(current_out || rescue) -> out[0..32]
3. XOF(hash || rescue) -> out[32..64]
```

KERNEL HANDOFF

Seed delivered to kernel RNG

```
seed_entropy()
collected = collect_boot_entropy(bs)
info.entropy64.copy_from_slice(&collected)
scrub(&mut collected) // Zero local copy
```

RTC TIMESTAMP

8-byte timestamp for misc use

```
get_rtc_timestamp() -> [u8; 8]
year[2] + month + day
hour + minute + second
milliseconds (from nanosecond)
```

CONSTANTS

```
HW_RNG_ITERATIONS = 32
TSC_JITTER_ROUNDS = 256
Output Size = 64 bytes
scrub() zeros sensitive data
```

ENTROPY FLOW SUMMARY

EFI_RNG (64B) + RDRAND/RDSEED (64B) + TSC Jitter (256 rounds) + RTC (16B)

-> BLAKE3 Accumulator -> Domain-Separated XOF -> 64-byte seed -> BootHandoffV1.rng.seed32 -> Kernel CSPRNG

NONOS Bootloader: Security Subsystem

TPM Measured Boot + UEFI Secure Boot + Self-Tests

TPM 2.0 MEASURED BOOT

src/security/tpm.rs

PCR Indices

```
pcr::BOOTLOADER = 8  
pcr::KERNEL = 9  
pcr::CAPSULE = 14
```

Standard TPM 2.0 PCRs

TCG2 Protocol

```
GUID: 607f766c-7455-  
42be-930b-e4d76db2720f  
  
EV_POST_CODE = 0x00000001  
HashLogExtendEvent offset: 0x18
```

extend_pcr_measurement()

1. Validate data not empty, pcr_index <= 23
2. SHA-256 hash of measurement data
3. locate_tcg2_protocol(bs)
4. extend_pcr_via_tcg2(tcg2, pcr, digest)
-> HashLogExtendEvent(flags=0x10, digest, event)
5. Return true on success, false if no TPM

UEFI SECURE BOOT CHECKS

src/security/check.rs

Platform Checks

```
check_secure_boot()  
UEFI SecureBoot variable  
  
check_platform_key()  
PK enrolled status  
  
check_signature_db()  
db/dbx databases
```

Hardware Checks

```
check.hardware_rng()  
RDRAND/RDSEED available  
  
check_measured_boot()  
TPM2 availability
```

All checks return bool

SecurityContext

```
pub struct SecurityContext {  
    secure_boot_enabled: bool,  
    measured_boot_available: bool,  
    hardware_rng_available: bool,  
    platform_key_enrolled: bool,  
}
```

MEASUREMENT FLOW

measure_boot_components()

```
PCR8: extend(bootloader_hash)          // Self measurement  
PCR9: extend(kernel_hash)             // Kernel binary  
PCR14: extend(capsule_hash)           // Full capsule  
  
Returns true only if ALL succeed
```

SIGNATURE VERIFICATION

verify_kernel_signature_advanced()

1. Extract signature from capsule
2. Hash kernel payload with BLAKE3
3. Ed25519 verify against keystore

Returns (valid: bool, keyid: Option<KeyId>)

Security Defense Layers

Layer 1: UEFI Secure Boot (firmware) | Layer 2: Ed25519 Signature (bootloader) | Layer 3: TPM Measurements (hardware) | Layer 4: ZK Attestation (cryptographic)

CRYPTOGRAPHIC SELF-TESTS

src/security/crypto.rs

blake3_selftest()

```
test_data = b"NONOS-BLAKE3-SELFTEST"  
hash1 = blake3::hash(test_data)  
hash2 = blake3::hash(test_data)  
return hash1 == hash2 // Determinism check
```

ed25519_selftest()

```
VerifyingKey::from_bytes(NONOS_SIGNING_KEY)  
return result.is_ok() // Key parseable
```

run_all_selftests()

```
blake3_selftest() && ed25519_selftest()  
Called during boot initialization
```

SECURITY POSTURE

assess_security_posture()

FULL: SecureBoot + TPM + HW-RNG + PK

PARTIAL: Some but not all features

MINIMAL: Basic signature verification only

NONE: No security features available

NONOS Bootloader: Network Boot

PXE/TFTP + HTTP Boot + DHCP Configuration

MODULE STRUCTURE

src/network/mod.rs

Core Modules:

```
dhcp_core // DHCP protocol  
tftp // TFTP client  
http // HTTP client  
pxe // PXE protocol
```

Support Modules:

```
retry // Retry logic  
signature // Verify downloads  
static_ip // Manual config  
diagnostics // Network debug
```

PXE / TFTP BOOT

Traditional Network Boot

PXE Functions:

```
is_pxe_available() -> bool  
get_tftp_server_from_dhcp() -> Ipv4  
get_boot_filename_from_dhcp() -> &str
```

TFTP Download:

```
tftp_download(server, filename, buf)  
-> TftpDownload { bytes_read, status }
```

Kernel Fetch:

```
fetch_kernel_via_pxe(bs, buf)  
-> Result<usize, NetworkError>
```

HTTP BOOT

Modern Network Boot

URL Parsing:

```
parse_url(url) -> (host, port, path)
```

HTTP Response:

```
HttpResponse {  
    status_code: u16,  
    content_length: usize,  
    body: &[u8]  
}
```

Kernel Fetch:

```
fetch_kernel_via_http(bs, url, buf)
```

DHCP CONFIGURATION

IP Address Assignment

DhcpConfig:

```
ip_address: Ipv4Addr  
subnet_mask: Ipv4Addr  
gateway: Ipv4Addr  
dns_server: Ipv4Addr
```

Functions:

```
configure_dhcp_via_pxe()  
get_dhcp_assigned_ip()  
is_dhcp_configured() -> bool
```

Static IP Alternative:

```
set_static_ip(StaticIpConfig)
```

RETRY CONFIGURATION

RetryConfig:

```
DEFAULT_MAX_RETRIES = 3  
DEFAULT_INITIAL_DELAY_US = 100_000  
DEFAULT_MAX_DELAY_US = 5_000_000  
NET_TIMEOUT_SECS = 30
```

Exponential backoff with jitter

DOWNLOAD VERIFICATION

verify_downloaded_kernel():

1. Compute BLAKE3 hash of kernel
2. Fetch .sig file from server
3. Ed25519 verify(kernel, sig)
4. Return verified kernel bytes

Same verification as local boot

NETWORK BOOT CONTEXT

NetworkBootContext:

```
boot_option: NetworkBootOption  
config: NetworkConfig  
kernel_url: Option<String>  
signature_url: Option<String>  
timeout_secs: u32
```

NetworkBootOption:

```
PxeTftp // Traditional PXE  
HttpBoot // UEFI HTTP Boot  
DhcpOnly // Config only  
Disabled // Local boot
```

NETWORK BOOT FLOW



MAX_KERNEL_SIZE = 64 MB | NET_MAX_RETRIES = 3 | NET_TIMEOUT_SECS = 30 | Supports UEFI HTTP Boot & Legacy PXE

NONOS Bootloader: Kernel Handoff

BootHandoffV1 Structure + Assembly Transfer Protocol

BootHandoffV1 Structure

src/handoff/types.rs - #[repr(C)]

Header (16 bytes)

```
magic: u32 = 0x4E4F4E4F // "NONO"  
version: u16 = 1  
size: u16 = size_of::<Self>()
```

Flags (u64) + Entry Point

```
WX (1<<0) | NXE (1<<1) | SMEP (1<<2) | SMAP (1<<3)  
UMIP (1<<4) | IDMAP_PRESERVED (1<<5) | FB_AVAILABLE (1<<6)  
ACPI_AVAILABLE (1<<7) | TPM_MEASURED (1<<8)  
SECURE_BOOT (1<<9) | ZK_ATTESTED (1<<10)  
entry_point: u64 - Kernel entry address
```

FramebufferInfo

```
ptr: u64 // FB base  
size: u64 // Total bytes  
width: u32 // Pixels  
height: u32 // Pixels  
stride: u32 // Bytes/row  
pixel_format: u32
```

Firmware Tables

```
acpi.rsdः: u64 // RSDP address  
smbios.entry: u64 // SMBIOS entry
```

Modules + Cmdline

```
modules: { ptr: u64, count: u32 } // Module array  
cmdline_ptr: u64 // Kernel command line
```

BootHandoffV1::is_valid()

```
magic == 0x4E4F4E4F && version == 1 && size == sizeof(Self)
```

SECURITY HANDOFF

Measurements

```
kernel_sha256: [u8; 32] // Kernel hash  
kernel_sig_ok: u8 // Signature valid  
secure_boot: u8 // UEFI SecureBoot  
zk_attestation_ok: u8 // ZK proof valid
```

ZkAttestation

```
verified: u8 // ZK verified flag  
flags: u8 // Attestation flags  
program_hash: [u8; 32] // Circuit identity  
capsule_commitment: [u8; 32]
```

RNG SEED

```
rng.seed32: [u8; 32] // Boot entropy  
From multi-source entropy collection
```

CryptoHandoff

Internal transfer structure

```
signature_valid: bool  
secure_boot: bool  
kernel_hash: [u8; 32]  
zk_attested: bool  
zk_program_hash, zk_capsule_commitment
```

ASSEMBLY KERNEL ENTRY

src/handoff/exit.rs - x86_64 Long Mode Transfer

jump_to_kernel()

```
// Calling convention: SystemV AMD64  
mov rdi, {handoff} // Arg 1: handoff ptr  
mov rsp, {stack} // New stack pointer  
xor rbp, rbp // Clear frame pointer  
jmp {entry} // Jump to kernel  
// Never returns (-> !)  
// Kernel: extern "C" fn(u64) -> !
```

Kernel Entry Signature

```
pub type KernelEntry =  
extern "C" fn(u64) -> !
```

Kernel receives:

RDI = BootHandoffV1 pointer
RSP = Fresh 16KB stack
RBP = 0 (no frame chain)
Long mode, paging enabled, no interrupts

HANDOFF_MAGIC = 0x4E4F4E4F ("NONO") | HANDOFF_VERSION = 1 | Full hardware state passed to kernel

NONOS Bootloader: ELF Loader Subsystem

ET_EXEC + ET_DYN Loading with Relocation Support

LOADER MODULES

src/loader/

Core:

core/load.rs // Main loader
core/alloc.rs // Page alloc

Validation:

validate/mod.rs // ELF checks
validate/program.rs

Support:

reloc.rs, image.rs, errors.rs

ELF VALIDATION

validate_elf() -> ValidationResult

Header Checks:

Magic: 0x7F 'E' 'L' 'F'
Class: ELFCLASS64
Machine: EM_X86_64
Type: ET_EXEC or ET_DYN

Segment Validation:

PT_LOAD segments only
Bounds checking, no overlaps

LOAD STRATEGIES

ET_EXEC (Static)

load_exec_kernel()
AllocateType::Address(base) - Fixed address
No relocations needed

ET_DYN (PIE)

load_dyn_kernel()
AllocateType::AnyPages - Relocatable
process_elf_relocations() applied

MEMORY ALLOCATION

UEFI Boot Services

Constants:

PAGE_SIZE = 4096
MAX_ALLOCS = 64

Functions:

record_alloc(addr, pages)
free_all(bs, allocations)

MemoryType::LOADER_DATA

KERNEL LOAD FLOW

1. Capsule Load

load_validated_capsule()
Verify sig, extract payload
or pass raw ELF

2. Validate ELF

validate_elf(payload)
Check header, segments
Calculate min/max addr

3. Allocate Pages

bs.allocate_pages()
ET_EXEC: fixed addr
ET_DYN: any pages

4. Copy Segments

copy_nonoverlapping()
Load PT_LOAD segments
Zero BSS regions

5. Relocate

process_elf_relocations()
ET_DYN only: apply
R_X86_64_RELATIVE

6. Return Image

KernelImage { ... }
address, size, entry
metadata, allocations

Segment Copy (for each PT_LOAD)

copy_nonoverlapping(payload[p_offset], dst, p_filesz)
write_bytes(dst + p_filesz, 0, p_memsz - p_filesz) // BSS

Relocation Processing

load_bias = base_phys - min_addr
R_X86_64_RELATIVE: *ptr += load_bias

Entry Point Validation

entry >= base && entry < base + total_bytes
Err(LoaderError::EntryNotInRange) on failure

KernelImage Structure

```
pub struct KernelImage {  
    address: usize, // Load address  
    size: usize, // Total size  
    entry_point: usize, // Entry address  
    metadata: CapsuleMetadata, allocations, alloc_count  
}
```

ValidationResult

```
pub struct ValidationResult {  
    elf: Elf64, // Parsed ELF  
    is_exec: bool, // ET_EXEC or ET_DYN  
    min_addr, max_addr: u64, // Address range  
    loads: [LoadSegment; 16], load_count: usize  
}
```

LoaderError

```
InvalidElfMagic, InvalidElfClass  
InvalidElfMachine, InvalidElfType  
TooManyLoadSegments, SegmentOverlap  
AllocationFailed, EntryNotInRange
```

Supports 64-bit ELF | ET_EXEC (fixed) + ET_DYN (PIE) | UEFI page allocation | R_X86_64_RELATIVE relocations

NØNOS Kernel Architecture

x86_64 Platform Abstraction Layer • ZeroState Privacy OS

NØNOS Kernel Subsystems

process/

Process Control Block
Context Switching
Uses: GDT, IDT, Syscall

memory/

Page Tables (4-Level)
Physical/Virtual Alloc
Uses: CPU, ACPI

sched/

Task Scheduling
Timer Interrupts
Uses: Time, IDT, APIC

interrupts/

IRQ Management
Handler Dispatch
Uses: IDT, APIC

syscall/

System Call Interface
User/Kernel Transition
Uses: arch::syscall

drivers/

Device Drivers
Hardware Abstraction
Uses: PCI, Port, DMA

fs/

Filesystem (VFS)
RAM-only Mode
Uses: Memory, Storage

security/

Capabilities
Audit, Hardening
ZeroState Security

crypto/

Ed25519, BLAKE3
ML-KEM, ML-DSA
Post-Quantum Ready

smp/

Multi-Processor
AP Startup, IPI
Uses: CPU, APIC, GDT

modules/

Signed Capsules
Sandbox, Registry
ZeroState Modules

ipc/

Inter-Process Comm
Shared Memory, Signals
Uses: Memory, Syscall

net/

Network Stack (smoltcp)
TCP/IP, Sockets
Uses: PCI, IRQ, DMA

ui/

User Interface
TUI, Graphics
Uses: VGA, Input

94 Modules | 15+ Subsystems | 256 IRQ Vectors | 6 IST Stacks | 256 Max CPUs

4

Core Path

Submodule

Kernel Layer

arch/x86_64/

cpu/

CPUID Detection
Feature Flags
MSR Access
TSC Operations
Per-CPU Data

gdt/

Segments
TSS Management
IST Stacks
Selectors
SYSCALL Setup

idt/

256 Vectors
Exception Handlers
IRQ Handlers
IST Assignment
Handler Registry

interrupt/

Local APIC
x2APIC Support
I/O APIC
Legacy 8259A PIC
IPI / MSI-X

time/

TSC (Primary)
HPET
PIT (8254)
RTC (CMOS)
Unified API

syscall/

SYSCALL/SYSRET
MSR Config
Handler Dispatch
Security Policies
Statistics

acpi/

RSDP/XSDT
MADT (CPUs)
FADT (Power)
HPET/MCFG
NUMA Tables

I/O & Boot

serial/ (16550A UART)
vga/ (Text Mode)
keyboard/ (PS/2, USB)
pci/ (Config Space)
boot/ | multiboot/ | uefi/

Hardware Layer

CPU x86_64

Memory

APIC/IOAPIC

PCI Bus

Timers

Serial/VGA

Keyboard

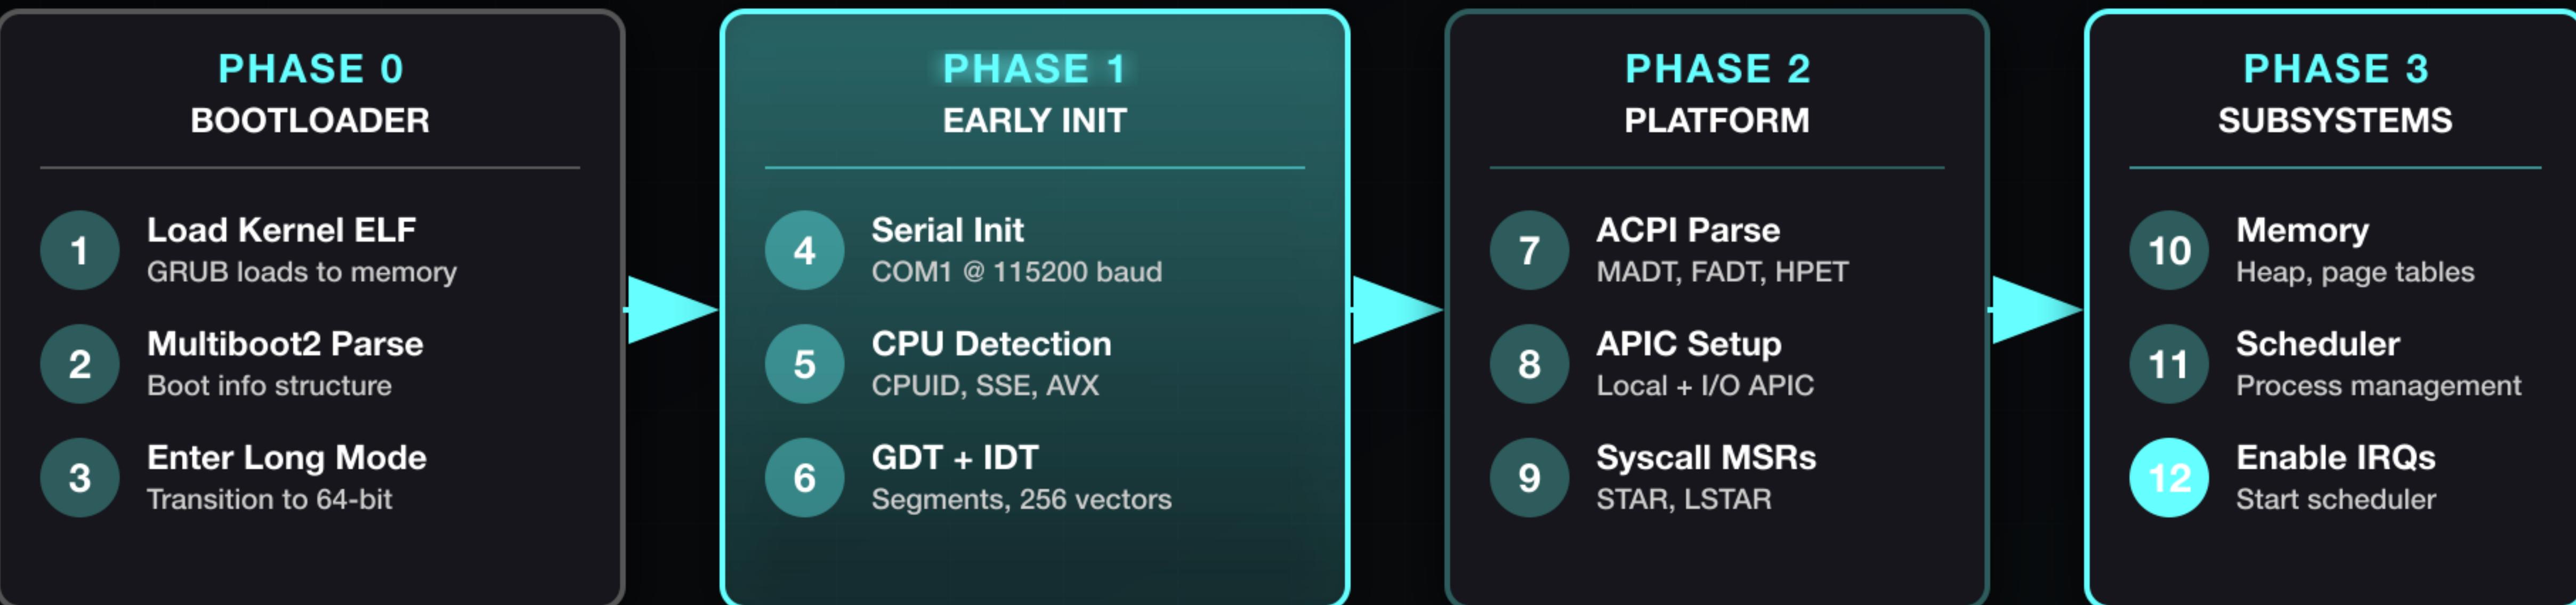
Storage

Network

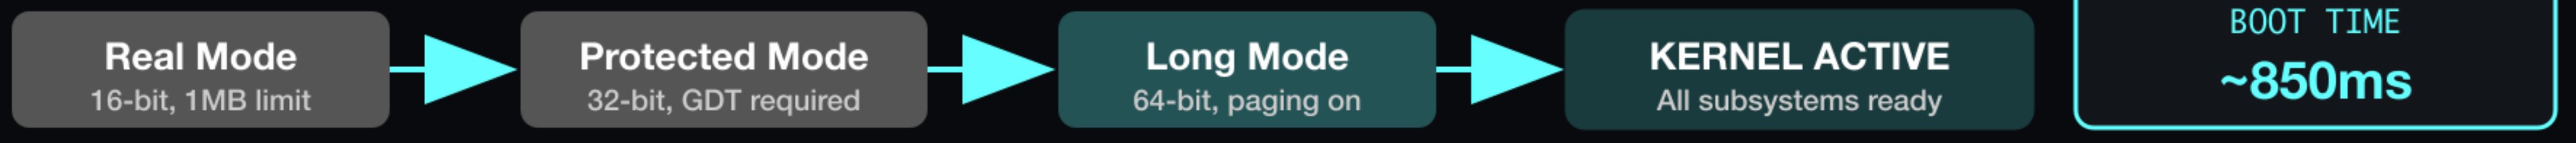
ACPI/UEFI

NØNOS Boot Sequence

x86_64 Kernel Initialization Flow



CPU Mode Transitions



SYSTEM READY - NØNOS Kernel v0.8.0

CPU Subsystem

arch/x86_64/cpu/ - Detection, Features & Management

CPU MODULE STRUCTURE

cpuid.rs
CPUID instruction wrapper

features.rs
100+ CPU feature flags

msr.rs
Model-Specific Registers

tsc.rs
Time Stamp Counter

topology.rs
Core/Thread detection

per_cpu.rs
Per-CPU data (256 max)

KEY MSR REGISTERS

0xC0000080 **IA32_EFER** Extended features

0xC0000081 **IA32_STAR** SYSCALL segments

0xC0000082 **IA32_LSTAR** SYSCALL entry

0xC0000100 **IA32_FS/GS** Segment bases

CPUID INSTRUCTION FLOW



CPU TOPOLOGY

PACKAGE 0

Core 0

Thread 0, 1

Core 1

Thread 0, 1

CACHE HIERARCHY

L1

32 KB
~4 cycles

L2

256 KB
~12 cycles

L3

8-16 MB
~40 cycles

CPU FEATURES (100+ FLAGS)

SIMD
SSE, SSE2, SSE3
AVX, AVX2, AVX-512
FMA, BMI1, BMI2
+20 more

SECURITY
AES-NI, SHA
RDRAND, RDSEED
SMEP, SMAP, CET
+15 more

TIME STAMP COUNTER

cpu::rdtsc() // Fast read (~20 cycles)

cpu::rdtscp() // Serialized + CPU ID

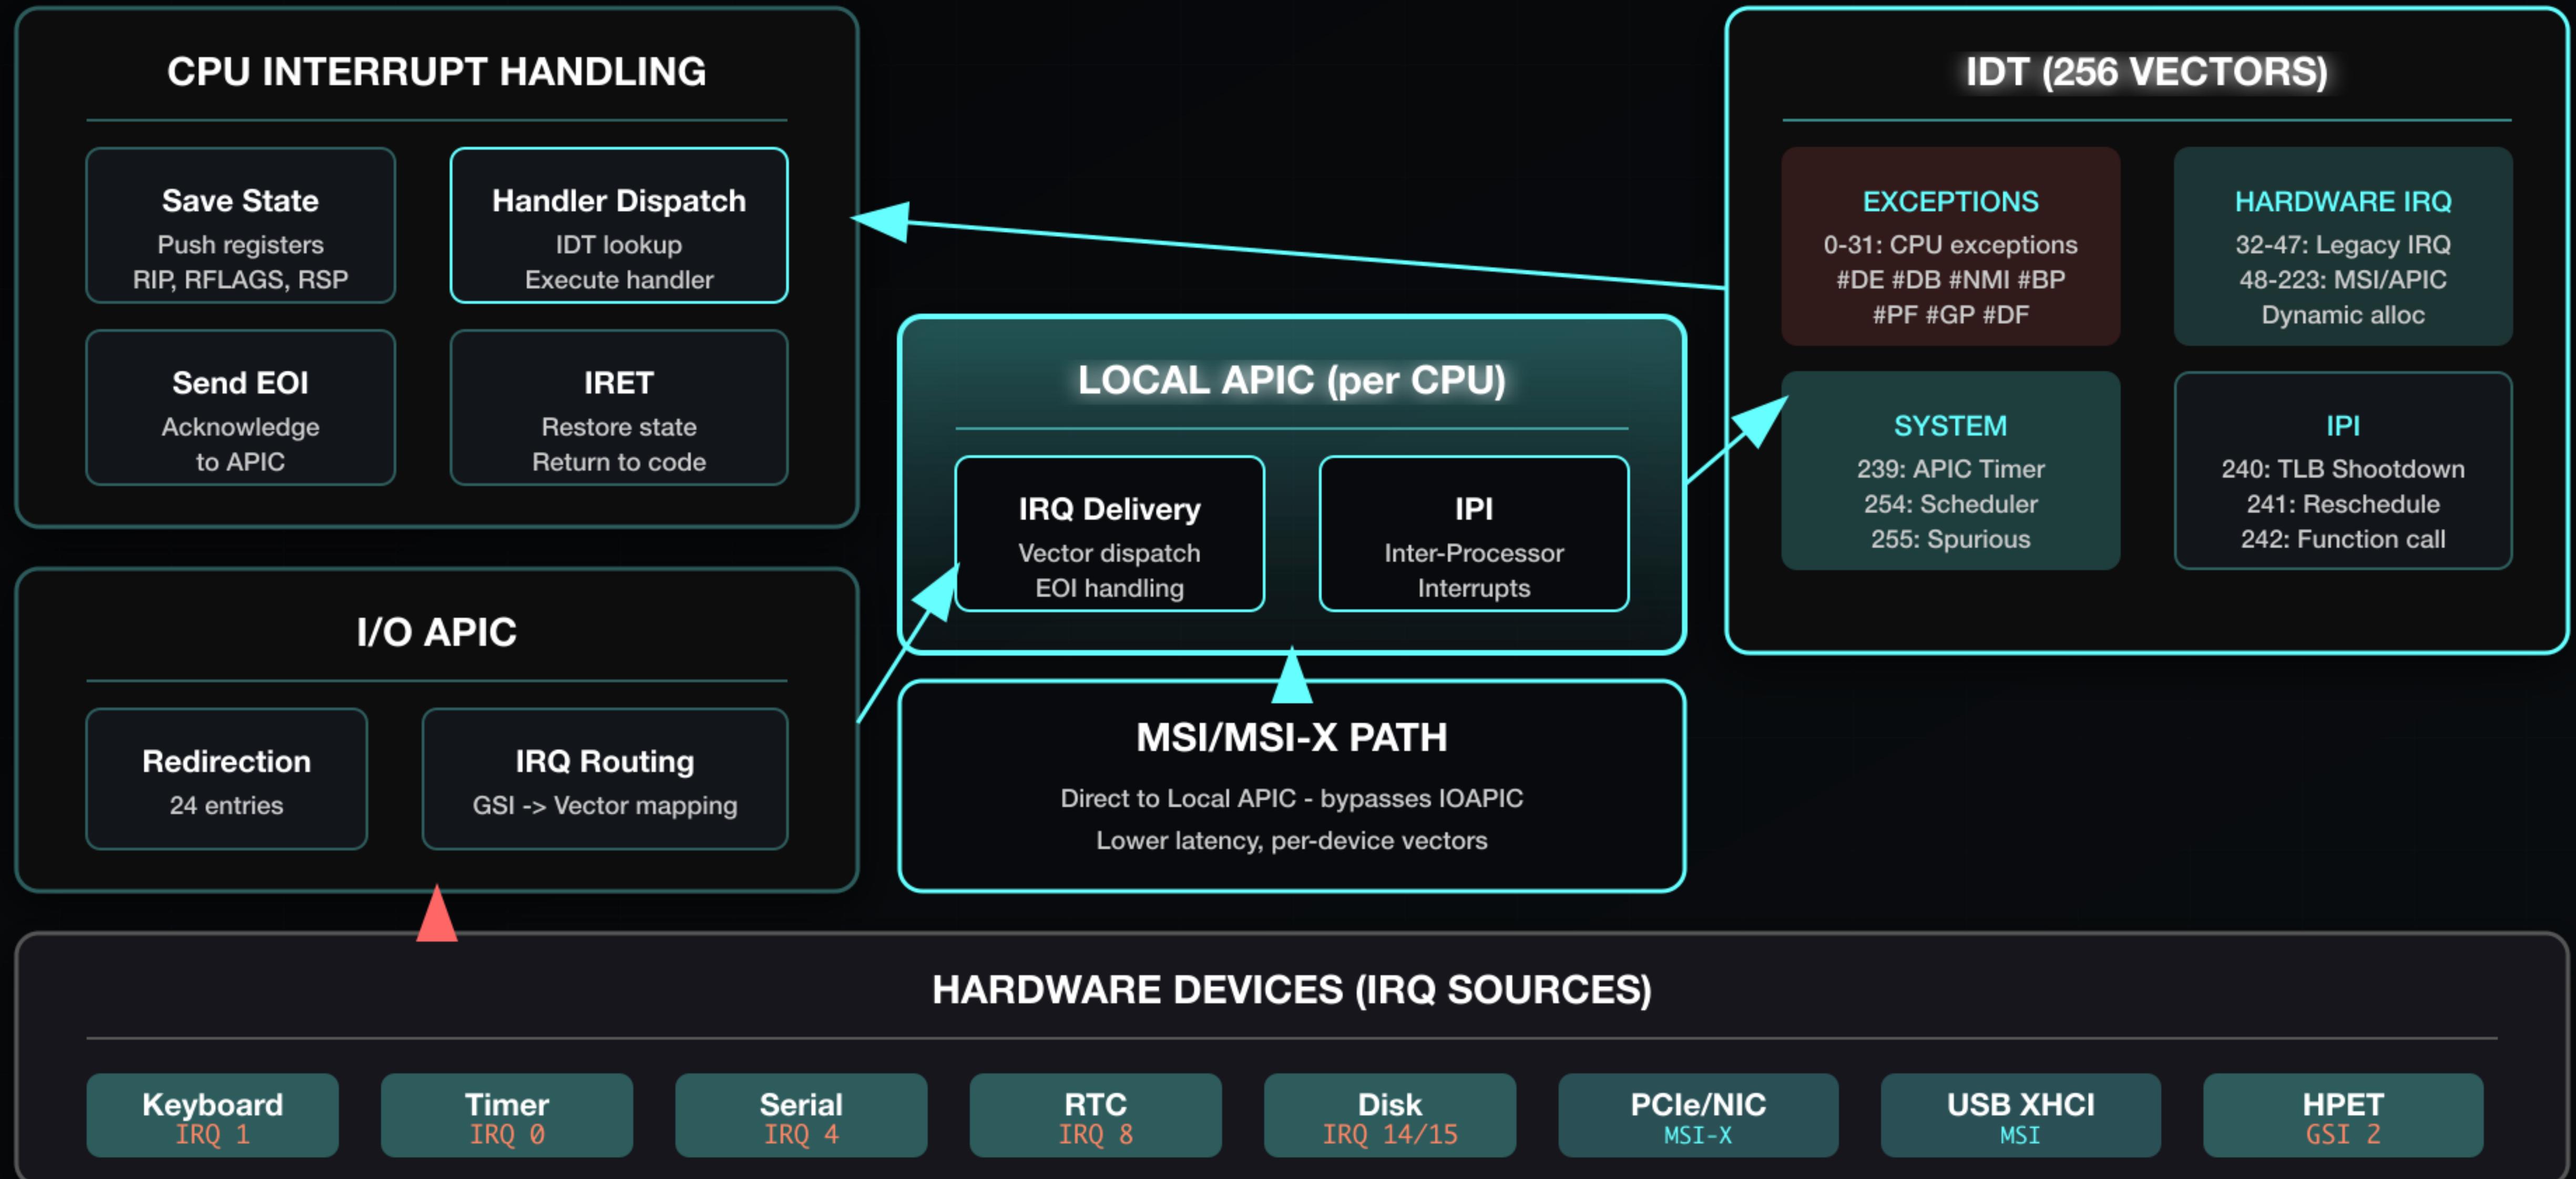
Invariant TSC: constant rate across P/C states

PER-CPU DATA

cpu_id: u32 // APIC ID
features: CpuFeatures // Cached
kernel_stack: u64 // via GS segment

Interrupt Architecture

arch/x86_64/interrupt/ - APIC, IOAPIC & IDT



Timer Subsystem

arch/x86_64/time/ - Multi-Source Timekeeping

CLOCK SOURCE HIERARCHY (Priority Order)

TSC

HPET

APIC Timer

PIT 8254

RTC

TSC (Time Stamp Counter)

PRIMARY - Highest Resolution

```
cpu::rdtsc() // ~20 cycles  
cpu::rdtscp() // Serialized + CPU ID
```

Resolution
~1 ns

Invariant TSC
Constant rate

HPET (High Precision)

SECONDARY - For Calibration

```
ACPI HPET table detection  
Memory-mapped registers
```

Resolution
~100 ns

Frequency
14.3 MHz+

APIC Timer

PER-CPU - Scheduling

```
Local APIC timer interrupt  
One-shot or periodic mode
```

Use Case
Preemption

Mode
TSC-deadline

UNIFIED TIME API

```
time::now() -> Instant  
time::uptime() -> Duration  
time::sleep(Duration)  
time::delay_us(u64)
```

```
time::tsc_frequency() -> u64  
time::calibrate_tsc()  
time::monotonic_ns() -> u64
```

FALLBACK SOURCES

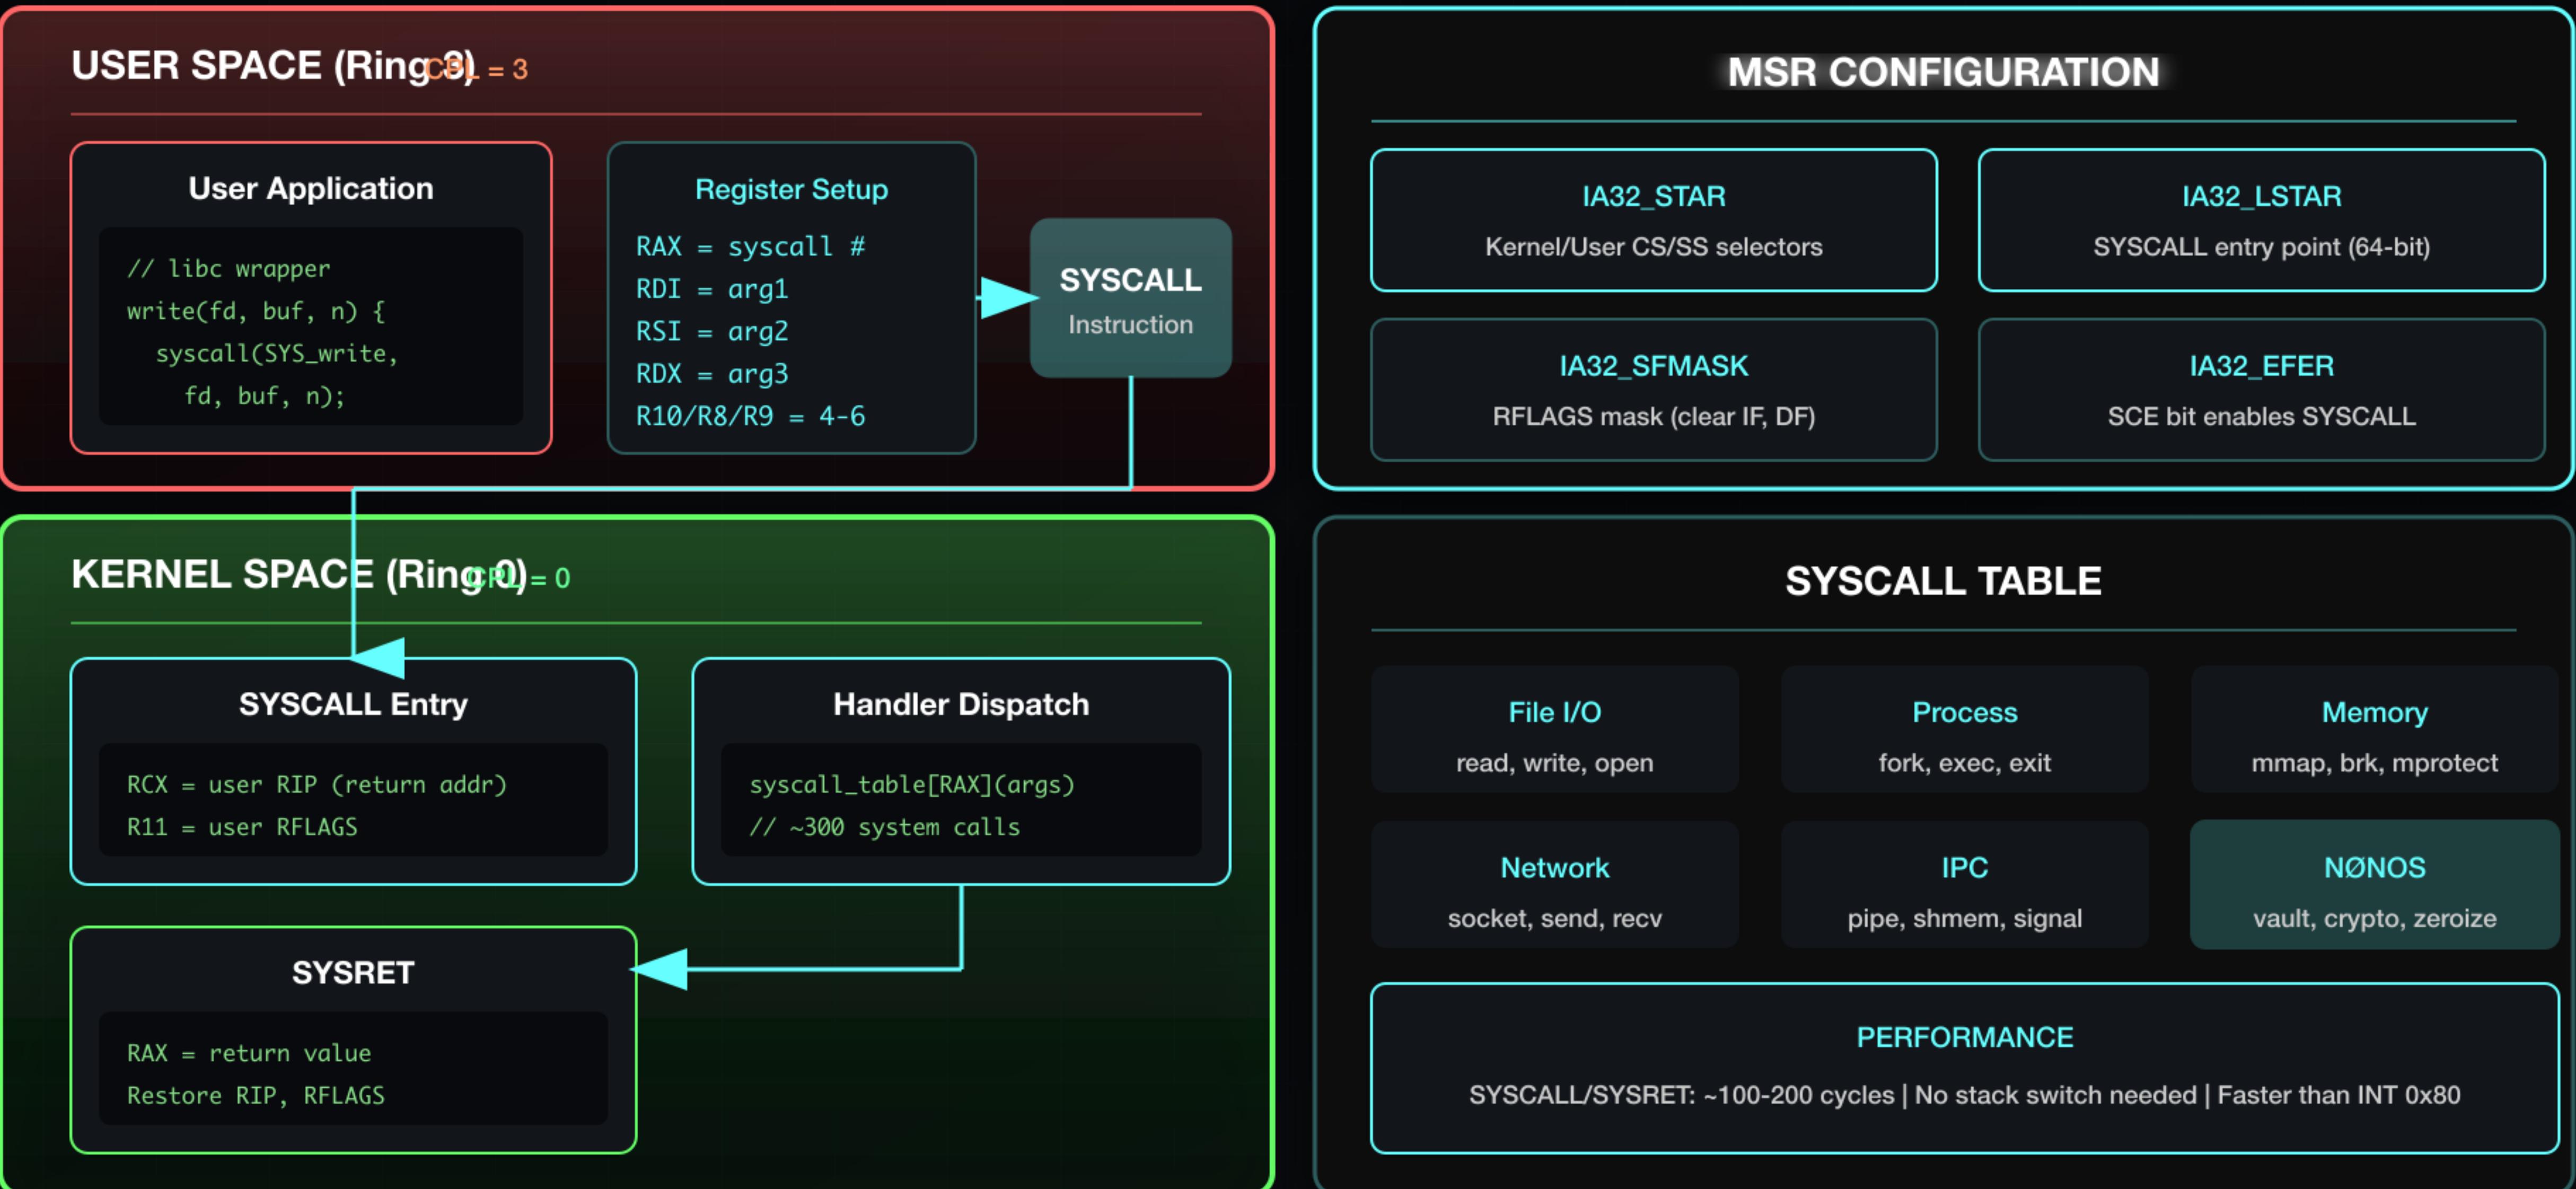
PIT 8254
Legacy timer
~840 ns
1.193 MHz

RTC (CMOS)
Wall clock time
1 second
Battery-backed

TSC Calib
PIT used to
calibrate TSC
frequency

SYSCALL / SYSRET Flow

arch/x86_64/syscall/ - User/Kernel Transition



Crypto Subsystem

Post-Quantum Ready Cryptographic Primitives

HASH FUNCTIONS (hash/)

BLAKE3 Primary	SHA-512 512-bit	SHA3 Keccak
SHA-256 unified/	HMAC MAC	HKDF KDF
SHA-1 Legacy only	Keccak256 / SHAKE128/256 Ethereum compatible	

SYMMETRIC (symmetric/)

ChaCha20-Poly1305 NØNOS Primary AEAD	AES-256-GCM AEAD
AES-128/256 Block cipher	Poly1305 MAC
core/aead.rs - Unified AEAD traits	

ASYMMETRIC (asymmetric/)

Ed25519 Signatures	X25519 Key exchange	P-256 NIST curve
secp256k1 Bitcoin/ETH	RSA Legacy	Curve25519 DH base
core/traits.rs - Kem, Sig, Ed25519Sig traits		

POST-QUANTUM (pqc/)

NIST PQC Standards Ready

ML-KEM (Kyber) kyber.rs - KEM	ML-DSA (Dilithium) dilithium.rs - Sig
SPHINCS+ sphincs/	McEliece mceliece/
quantum.rs - Hybrid PQ utilities	

ZERO-KNOWLEDGE (zk/)

Groth16 zkSNARK	Halo2 Recursive	nonos_zk Attestation
zk_kernel/ - Native ZK Proofs		
Pedersen Schnorr Sigma Range Equality Membership PLONK syscall_zk_verify, syscall_zk_commit, syscall_zk_prove_*		

APPLICATION (application/)

vault.rs Secure key storage	nonos_signing.rs Module signatures
ethereum/ keccak256, ecrecover	certification/ X.509 parsing

CORE API (core/)

api.rs init_crypto_subsystem generate_keypair, sign, verify	aead.rs Aead trait, aead_wrap/unwrap
syscall.rs syscall_blake3_hash, ...	traits.rs Kem, Sig, KyberKem, ...

UTILITIES (util/)

bigint/ Arbitrary precision math	entropy.rs RDRAND/RDSEED
constant_time/ Side-channel safe	rng/ get_random_bytes

SECURITY PROPERTIES

ZeroState Keys RAM-only Zeroize on drop	Side-Channel Constant-time ops No timing leaks	Quantum Ready ML-KEM + ML-DSA Hybrid mode
---	--	---

Memory safety via Rust | No unsafe crypto | NIST compliance | Test vectors validated

NØNOS DEFAULT ALGORITHM SELECTION

Symmetric AEAD ChaCha20-Poly1305 Primary for all encryption	Signatures Ed25519 + ML-DSA (hybrid) Module signing, auth	Key Exchange X25519 + ML-KEM (hybrid) Forward secrecy
---	---	---

Feature flags: mlkem512/768/1024, mldsa2/3/5, zk-halo2, zk-groth16

Hash BLAKE3 (fast) / SHA-512 Integrity, KDF input

Driver Subsystem

drivers/ - 188 files, 23 modules - Hardware Abstract

Common Pattern

```
static CONTROLLER: spin::Once<T> = Once::new();  
Single initialization, lock-free access after init
```

DRIVER PUBLIC API (mod.rs exports)

Initialization

```
init_all_drivers() init_nvme()  
init_ahci() init_xhci() init_tpm()
```

Controller Access

```
get_nvme_controller()  
get_ahci_controller() get_xhci_*
```

Network

```
wifi_connect() wifi_scan()  
get_virtio_net_device()
```

TPM / Security

```
measure_module() read_pcr()  
tpm_get_random_bytes()
```

Display / Audio

```
init_gpu() with_gpu_driver()  
get_audio_controller()
```

Stats / Info

```
get_hardware_stats()  
get_all_devices() DeviceInfo
```

STORAGE DRIVERS

NVMe

```
NvmeController  
Admin/I/O queues  
PRP/SGL lists  
Namespace mgmt
```

AHCI

```
AhciController  
FIS-based cmds  
Command slots  
Secure erase
```

Controller Submodules

```
controller/core.rs init.rs io.rs  
controller/admin.rs commands.rs  
dma.rs queue.rs security.rs stats.rs  
types.rs constants.rs error.rs
```

NETWORK DRIVERS

VirtIO-Net

```
VirtQueue  
modern_regs  
smoltcp  
PacketBuf
```

E1000

```
RX/TX rings  
EEPROM  
PHY ctrl  
checksum
```

Intel WiFi

```
firmware  
scan/conn  
tx/rx DMA  
PCIe regs
```

Network Stack Integration

```
VirtioNetInterface VirtioSmolBridge  
interface.rs validation.rs rate_limiter  
RTL8139 (legacy) network/stack.rs  
EtherType parsing MAC address mgmt
```

USB / AUDIO / DISPLAY

xHCI

```
TRB rings  
Event/Cmd  
DeviceCtx  
USB 3.x
```

HD Audio

```
BDL DMA  
Codec verb  
Streams  
PCM fmt
```

GPU

```
VBE modes  
Framebuf  
Backbuf  
32bpp
```

Controller Details

```
CommandRing EventRing TransferRing  
SlotContext EpContext InputContext  
VGA text Console framebuf Keyboard  
AudioFormat StreamState DisplayMode
```

SECURITY / TPM

TPM 2.0

```
extend_pcr_sha256()  
measure_component()  
verify_boot_chain()  
PCR 0-23 banks
```

security/

```
validate_dma_buffer()  
validate_mmio_region()  
validate_lba_range()  
RateLimiter
```

Security Primitives

```
safe_mmio_read32/write32  
is_config_write_allowed()  
DMA bounds LBA partition PRP validation  
Boot measurement log Quote generation
```

PCI/PCIe BUS LAYER

PciManager

```
scan_and_collect()  
find_device_by_class()  
find_device_by_id()
```

ConfigSpace

```
read32/write32  
BAR decoding  
0xCF8/0xCFC ports
```

MSI / MSI-X

```
configure_msi()  
configure_msix()  
Vector allocation
```

Capabilities

```
enumerate_capabilities()  
get_PCIE_info()  
PM, AER, SR-IOV
```

SecurityPolicy

```
validate_config_write()  
approve_bus_master()  
Allow/Block lists
```

DmaEngine

```
DmaDescriptor  
scatter-gather  
64-bit addressing
```

PciStats

```
config reads  
config writes  
violations
```

HARDWARE LAYER

NVMe SSD

PCIe 4.0 x4

SATA HDD

AHCI 6Gbps

Intel NIC

E1000/E1000E

VirtIO NIC

QEMU/KVM

Intel WiFi

AX200/AX210

USB 3.x

xHCI

HD Audio

Intel HDA

GPU

Bochs VBE

TPM 2.0

MMIO/LPC

PS/2

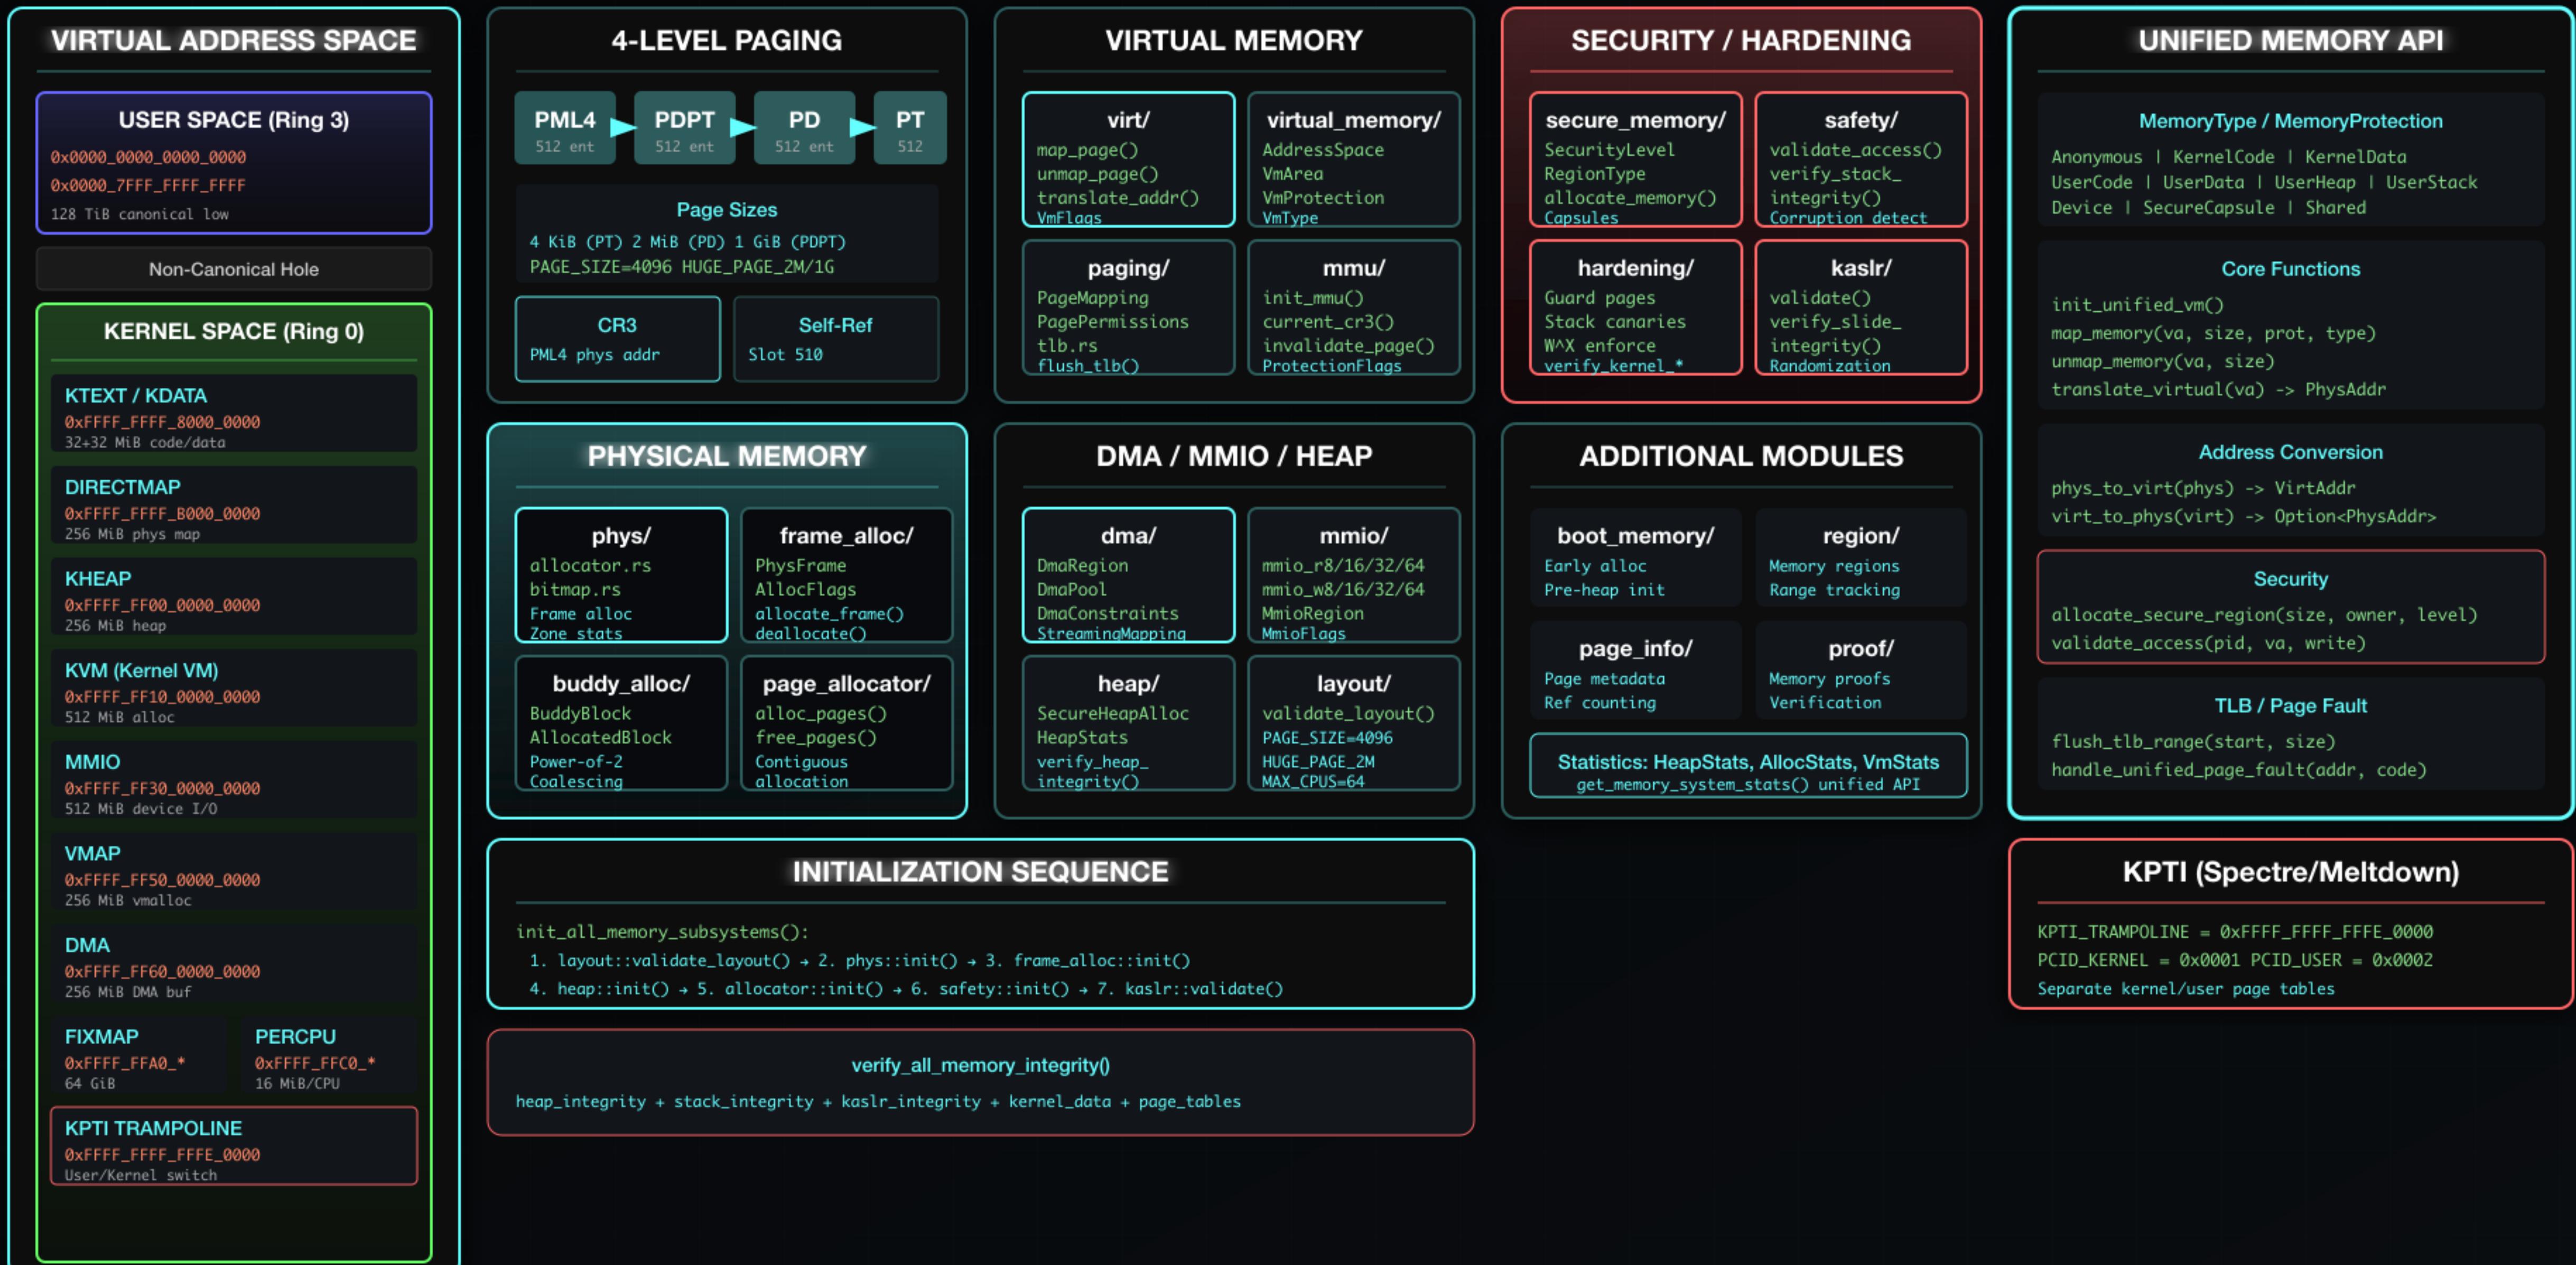
8042 KBC

RTL8139

Legacy NIC

Memory Subsystem

memory/ - 21 modules - Physical, Virtual, DMA, Security



ELF Subsystem

elf/ - 21 modules - Executable and Linkable Format Parser & Loader

ELF FILE STRUCTURE

ElfHeader (64 bytes)

0x7F 'E' 'L' 'F'
e_type: ET_EXEC/DYN/REL
e_machine: EM_X86_64
e_entry: entry point
e_phoff, e_shoff

ProgramHeader[]

PT_LOAD Loadable segment
PT_DYNAMIC Dynamic linking
PT_INTERP Interpreter path
PT_TLS Thread-local
PT_GNU_RELRO Read-only reloc
PT_GNU_STACK Stack perms

SectionHeader[]

SHT_PROGBITS .text/.data
SHT_SYMTAB Symbol table
SHT_STRTAB String table
SHT_REL A Relocations
SHT_DYNAMIC Dynamic info

DynamicEntry[]

DT_NEEDED Dependencies
DT_SYMTAB Symbol table
DT_STRTAB String table
DT_REL A Relocations
DT_INIT/FINI Init/Fini funcs

Symbol (24 bytes)

st_name, st_info
st_value, st_size
STB_GLOBAL/LOCAL/WEAK

RelaEntry (24 bytes)

r_offset, r_info
r_addend
type | sym << 32

CORE MODULES

types/

ElfHeader
ProgramHeader
SectionHeader
Symbol, RelaEntry

loader/

ElfLoader
ElfImage
LoadedSegment
DynamicInfo

reloc/

process_relocations()
RelocationContext
apply.rs, utils.rs
x86_64_relocs

symbol/

SymbolLookup
SymbolResolver
ResolvedSymbol
Name_resolution

HASH / GOT / LINKING

hash/

gnu_hash()
sysv_hash()
DualHashLookup

got/

GlobalOffsetTable
GotEntry
PLT/GOT_relocs

dynlink/

DynLinkInfo
Dynamic_linking_info_parsing

libmgr/

LibraryManager
LinkMap

interpreter/

InterpreterInfo
ld-linux.so

aslr/

AslrManager
Randomization

x86_64 RELOCATION TYPES

R_X86_64_64 S + A (absolute 64-bit)
R_X86_64_PC32 S + A - P (PC-relative)
R_X86_64_GLOB_DAT S (GOT entry)
R_X86_64_JUMP_SLOT S (PLT entry)
R_X86_64_RELATIVE B + A (PIE/DSO)
R_X86_64_COPY Copy symbol
R_X86_64_PLT32 L + A - P
R_X86_64_GOTPCREL G + GOT + A - P

TLS Relocations:

R_X86_64_DTPMOD64 TLS module ID
R_X86_64_DTPOFF64 TLS offset
R_X86_64_TPPOFF64 TP offset

ADDITIONAL

minimal/ validate_elf

cache/
ImageCache

errors/ ElfError

constants
ELF_MAGIC

PUBLIC API (mod.rs)

load_elf_executable()
init_elf_loader()
get_elf_loader()
validate_elf()
validate_elf_detailed()
create_process()
create_process_with_args()
process_relocations()
setup_user_stack()

Re-exports:

ElfHeader, ProgramHeader
Symbol, RelaEntry
ElfImage, LoadedSegment

ELF TYPES & MACHINES

e_type:

ET_EXEC (2) Executable
ET_DYN (3) Shared/PIE
ET_REL (1) Relocatable

e_machine:

EM_X86_64 (62)

ELF LOADING FLOW

1. Validate

validate_elf()

2. Parse

ElfLoader::load()

3. Map

PT_LOAD segments

4. Relocate

process_relocs()

5. Setup Stack

setup_user_stack

6. Execute

jump to e_entry

User Stack Layout (setup by auxv/stack)

| argc | argv[] | NULL | envp[] | NULL | auxv[] | AT_NULL |